

CS260R Reinforcement Learning

Assignment 0: Jupyter Notebook usage and assignment submission workflow

CS260R 2023Fall: Reinforcement Learning. Department of Computer Science at University of California, Los Angeles. Course Instructor: Professor Bolei ZHOU. Assignment author: Zhenghao PENG, Yiran WANG.

You are asked to finish four tasks:

1. Fill in your name and University ID in the next cell.
2. Install pytorch and finish the Kindergarten Pytorch section.
3. Run all cells and save this notebook **as a PDF file**.
4. Compress this folder `assignment0` **as a ZIP file** and submit **the PDF file and the ZIP file separately as two files** in BruinLearn.

```
In [14]: # TODO: Fill your name and UID here
my_name = "Shaira"
my_student_id = "506302126"
```

```
In [15]: # Run this cell without modification

text = "Oh, I finished this assignment! I am {} ({}).".format(my_name, my_student_id)
print(text)
with open("{}_txt".format(text), "w") as f:
    f.write(text)
```

Oh, I finished this assignment! I am Shaira (506302126)

Kindergarten Pytorch

1. Please install pytorch in your virtual environment following the instruction:

<https://pytorch.org/get-started/locally/>.

```
pip install torch torchvision
```

2. If you are not familiar with Pytorch, please go through [the tutorial in official website](#) until you can understand the [quick start tutorial](#).
3. The following code is copied from the [quick start tutorial](#), please solve all `TODO` s and print the result in the cells before generating the PDF file.

Prepare data

```
In [16]: import torch
from torch import nn
from torch.utils.data import DataLoader
from torchvision import datasets
from torchvision.transforms import ToTensor

# Download training data from open datasets.
training_data = datasets.FashionMNIST(
    root="data",
    train=True,
    download=True,
    transform=ToTensor(),
)

# Download test data from open datasets.
test_data = datasets.FashionMNIST(
    root="data",
    train=False,
    download=True,
    transform=ToTensor(),
)

batch_size = 64

# Create data loaders.
train_dataloader = DataLoader(training_data, batch_size=batch_size)
test_dataloader = DataLoader(test_data, batch_size=batch_size)
```

Define model

```
In [17]: # Get cpu, gpu or mps device for training.
device = (
    "cuda"
    if torch.cuda.is_available()
    else "mps"
    if torch.backends.mps.is_available()
    else "cpu"
)
print(f"Using {device} device")

# Define model
class NeuralNetwork(nn.Module):
    def __init__(self):
        super().__init__()
        self.flatten = nn.Flatten()

    # TODO: Define the self.linear_relu_stack by uncommenting next few lines
    # and understand what they mean
    self.linear_relu_stack = nn.Sequential(
        nn.Linear(28*28, 512),
        nn.ReLU(),
        nn.Linear(512, 512),
        nn.ReLU(),
        nn.Linear(512, 10)
```

```

    )

    def forward(self, x):
        x = self.flatten(x)
        logits = self.linear_relu_stack(x)
        return logits

model = NeuralNetwork().to(device)
print(model)

```

Using cpu device

```

NeuralNetwork(
  (flatten): Flatten(start_dim=1, end_dim=-1)
  (linear_relu_stack): Sequential(
    (0): Linear(in_features=784, out_features=512, bias=True)
    (1): ReLU()
    (2): Linear(in_features=512, out_features=512, bias=True)
    (3): ReLU()
    (4): Linear(in_features=512, out_features=10, bias=True)
  )
)

```

Define training and test pipelines

```

In [18]: loss_fn = nn.CrossEntropyLoss()
optimizer = torch.optim.SGD(model.parameters(), lr=1e-3)

def train(dataloader, model, loss_fn, optimizer):
    size = len(dataloader.dataset)
    model.train()
    for batch, (X, y) in enumerate(dataloader):
        X, y = X.to(device), y.to(device)

        # Compute prediction error
        pred = model(X)
        loss = loss_fn(pred, y)

        # Backpropagation

        # TODO: Uncomment next three lines and understand what they mean
        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

        if batch % 100 == 0:
            loss, current = loss.item(), (batch + 1) * len(X)
            print(f"loss: {loss:>7f} [{current:>5d}/{size:>5d}]")

def test(dataloader, model, loss_fn):
    size = len(dataloader.dataset)
    num_batches = len(dataloader)
    model.eval()
    test_loss, correct = 0, 0
    with torch.no_grad():
        for X, y in dataloader:

```

```
X, y = X.to(device), y.to(device)
pred = model(X)
test_loss += loss_fn(pred, y).item()

# TODO: Uncomment next line and understand what it means
correct += (pred.argmax(1) == y).type(torch.float).sum().item()

test_loss /= num_batches
correct /= size
print(f"Test Error: \n Accuracy: {(100*correct):>0.1f}%, Avg loss: {test_loss:>
```

Run the training and test pipelines

```
In [19]: epochs = 5
for t in range(epochs):
    print(f"Epoch {t+1}\n-----")
    train(train_dataloader, model, loss_fn, optimizer)
    test(test_dataloader, model, loss_fn)
print("Done!")
```

Epoch 1

```
-----  
loss: 2.326136 [ 64/60000]  
loss: 2.312474 [ 6464/60000]  
loss: 2.295205 [12864/60000]  
loss: 2.270647 [19264/60000]  
loss: 2.264616 [25664/60000]  
loss: 2.229636 [32064/60000]  
loss: 2.240818 [38464/60000]  
loss: 2.206999 [44864/60000]  
loss: 2.204479 [51264/60000]  
loss: 2.170034 [57664/60000]  
Test Error:  
Accuracy: 32.6%, Avg loss: 2.169871
```

Epoch 2

```
-----  
loss: 2.189871 [ 64/60000]  
loss: 2.181496 [ 6464/60000]  
loss: 2.127673 [12864/60000]  
loss: 2.126892 [19264/60000]  
loss: 2.091767 [25664/60000]  
loss: 2.034718 [32064/60000]  
loss: 2.059607 [38464/60000]  
loss: 1.985725 [44864/60000]  
loss: 1.983254 [51264/60000]  
loss: 1.918121 [57664/60000]  
Test Error:  
Accuracy: 59.1%, Avg loss: 1.915930
```

Epoch 3

```
-----  
loss: 1.954952 [ 64/60000]  
loss: 1.926252 [ 6464/60000]  
loss: 1.814991 [12864/60000]  
loss: 1.834932 [19264/60000]  
loss: 1.738485 [25664/60000]  
loss: 1.696647 [32064/60000]  
loss: 1.709980 [38464/60000]  
loss: 1.612477 [44864/60000]  
loss: 1.623685 [51264/60000]  
loss: 1.525933 [57664/60000]  
Test Error:  
Accuracy: 61.6%, Avg loss: 1.539463
```

Epoch 4

```
-----  
loss: 1.611671 [ 64/60000]  
loss: 1.571063 [ 6464/60000]  
loss: 1.424820 [12864/60000]  
loss: 1.480913 [19264/60000]  
loss: 1.367219 [25664/60000]  
loss: 1.367302 [32064/60000]  
loss: 1.377346 [38464/60000]  
loss: 1.300812 [44864/60000]  
loss: 1.325458 [51264/60000]
```

```

loss: 1.233728 [57664/60000]
Test Error:
Accuracy: 64.0%, Avg loss: 1.257738

```

Epoch 5

```

-----
loss: 1.339523 [ 64/60000]
loss: 1.316521 [ 6464/60000]
loss: 1.153039 [12864/60000]
loss: 1.247074 [19264/60000]
loss: 1.129029 [25664/60000]
loss: 1.155618 [32064/60000]
loss: 1.174312 [38464/60000]
loss: 1.111033 [44864/60000]
loss: 1.140478 [51264/60000]
loss: 1.064728 [57664/60000]
Test Error:
Accuracy: 65.0%, Avg loss: 1.085212

```

Done!

Save model

```

In [20]: torch.save(model.state_dict(), "model.pth")
         print("Saved PyTorch Model State to model.pth")

```

Saved PyTorch Model State to model.pth

Load model and run the inference

```

In [21]: model = NeuralNetwork().to(device)
         model.load_state_dict(torch.load("model.pth"))

```

Out[21]: <All keys matched successfully>

```

In [22]: classes = [
         "T-shirt/top",
         "Trouser",
         "Pullover",
         "Dress",
         "Coat",
         "Sandal",
         "Shirt",
         "Sneaker",
         "Bag",
         "Ankle boot",
         ]

         model.eval()
         x, y = test_data[0][0], test_data[0][1]
         with torch.no_grad():
             x = x.to(device)
             pred = model(x)

```

```
predicted, actual = classes[pred[0].argmax(0)], classes[y]  
print(f'Predicted: "{predicted}", Actual: "{actual}"')
```

Predicted: "Ankle boot", Actual: "Ankle boot"